

Question 1 (Swendsen-Wang):

(a, b): To understand SW algorithm, some background information is needed. There is a fundamental theorem [Robert and Casella, 2004, p47] that underlies the slice sampler and also SW algorithm. Assuming f is the pdf that we want to draw samples from, $f(x)$ can be written as

$$f(x) = \int_0^{f(x)} 1 du$$

$f(x)$ can be seen as the marginal distribution of joint variables (x, u)

$$(x, u) \sim \mathcal{U}\{(x, u) : 0 < u < f(x)\}, \quad (1)$$

where \mathcal{U} is the uniform distribution, and u is usually named as *auxiliary variable*. Thus, instead of drawing samples from $f(x)$ directly (which might be difficult), we can draw samples (x, u) from their uniform joint distribution on the constrained set $\{(x, u) : 0 < u < f(x)\}$. Once we have the samples, we can just discard u and x will be in original target distribution. This is the basic idea of slice sampler.

In slice sampler, we can generate a Markov chain with the stationary distribution equal to the joint uniform distribution of (1). We can generate x and u from their conditional distribution iteratively in a random-walk style: 1) generate u from $\mathcal{U}(\{u : u \leq f(x)\})$, and 2) given the new sample u , generate x from $\mathcal{U}(\{x : f(x) \leq u\})$. Robert and Casella [2004, p323] proves this Markov chain's stationary distribution is actually (1).

When $f(x)$ is complex function, finding the set of x such that $f(x) \leq u$ can be difficult (step 2 in above procedure). Such case can happen when the x is of large dimension (as in our fMRI study). The general slice sampler solves this problem by using multiple slices. In short, $f(x)$ can be factorized into the products of $f_i(x)$, and each f_i is associated with an auxiliary variable u_i . In this way, the support of the conditional distribution $p(x|u)$ in step 2 can be easily found. The SW sampler can be seen as one of such general slice samplers.

Here is the settings of the SW algorithm. Suppose x is the set of random variables defined on the lattice, and the probability of x is defined as

$$p(x) \propto \exp(-U(x)), \quad U(x) = - \sum_{(i,j)} \psi \mathbf{1}_{x_i=x_j}. \quad (2)$$

The indicator function $\mathbf{1}_{x_i=x_j}$ takes value 1 when the variables at two spatial adjacent sites are equal $x_i = x_j$. This is a standard Potts model. To sample from it using SW algorithm, we introduce a set of augmented variable u just like what we do in slice sampler. u corresponds to the bonds between spatially adjacent nodes. For each u_{ij} between x_i and x_j , there are two states *open* or *close* denoted by $u_{ij} = 1$ or $u_{ij} = 0$. The SW algorithm iterates between two steps

- Given x , set $u_{ij} = 0$ if $x_i \neq x_j$. When $x_i = x_j$, set $u_{ij} = 1$ with probability $1 - \exp(-\psi)$, and set $u_{ij} = 0$ with probability $\exp(-\psi)$. After this step, we have multiple connected components, each being a subset of the nodes on the graph.

- Given u , set all the nodes in a randomly chosen cluster (i.e. a connect component) with same label. The label is drawn from a uniform distribution.

Like slice sampler, the stationary distribution of this Markov chain is the joint distribution of u and x , which is again a uniform distribution [Rubinstein and Kroese, 2008, ch. 6.5.2]. And it can be proved [Winkler, 2003, ch. 8.2] that the marginal distribution $p(x)$ is exactly (2). So the joint model is consistent with the original marginal distribution. If we sum out x and get the marginal distribution of the augmented variable $p(u)$, we have a distribution called *random cluster model* [Grimmett, 2006, ch. 8]. After the sampling, we get samples of (x, u) , and we just ignore the augmented variable u , and x will be a sample from the original Potts model.

The SW sampling is more efficient than Gibbs sampling because at each step, it changes the labels of the whole cluster, instead of a single site. Even in low temperature, the sampler flip the labels for larger clusters. The mixing time of the sampling is polynomial in regular lattice.

[Barbu and Zhu, 2005] talked about the convergence rate of Swendsen-Wang (S-W) algorithm on a Potts model. They cited Huber [2003] who developed a new bounding chain algorithm that can diagnose the convergence of Swendsen-Wang sampling. The number of steps to reach perfect sampling (which means convergence to stationary distribution) is in the order of $\mathcal{O}(\log |E|)$, where E is the set of edges. This running time applies when the temperature is far below or far above critical temperature. Cooper and Frieze [1999] shows the mixing time (or convergence time) is polynomial if the number of neighbors of each node does not increase with $|V|$, the size of the nodes. This is good for regular lattice where the number of adjacent nodes is fixed, regardless of image size. Compared with the super-exponential rate of increase for the iteration number in standard Gibbs sampling (see question 5 of this report), S-W algorithm is a big improvement for convergence rate. One thing that need note is these theoretical analysis is for the cases without external fields (data likelihood term).

(c,d): The original SW sampling only applies to Ising and Potts models, and generalized SW sampling can be applied when there is data likelihood and we seek sampling from posterior probability. Another disadvantage of original SW is it is slow when sampling the labels for a cluster in step 2, since it does not use the observed data. The last contribution of generalized SW is the adaptive increasing or decreasing the number of labels.

There are two significant changes from standard SW to generalized SW. First, the probability of turning on the edges (bonds) at step 1 is changed from $q_0 = 1 - \exp(-\psi)$ to $q_e = -g(h_i, h_j)$, where h is the observed data (or features). $g(h_i, h_j)$ would takes larger value when the observed data at i and j are similar. The similarity is represented by the KL divergence in Barbu and Zhu [2005], but can be defined differently in other applications. Second, The sampling of labels in step 2 could be have acceptance rate smaller than one, instead of the 100% acceptance in original SW algorithm. the acceptance probability to move to new labels also depends on posterior probability given the observed data, as shown in theorem 2 of Barbu and Zhu.

The third version SWC-3 of the generalized SW replace the Metropolis-Hasting sampling

in step 2 with a Gibbs sampler and achieve the acceptance of 1. The Gibbs sampler draw labels from the posterior probability given the data.

Relating to the fMRI application, there are two issue to address in order to use SW sampling. First we need to define a function $g(h_i, h_j)$ to replace the KL divergence in eq (12) of Barbu and Zhu. The function will be plugged into the acceptance probability when sampling the augmented variable u (edge variables). One straightforward solution is to use the correlation between two time courses. An edge will be open with larger probability if two voxels connecting the edge has higher correlation. However, there is no justification of such method. Actually, more work need to be done to find the relationship between the acceptance probability q_e and the marginal probability $p(label|data)$.

Another issue is the sampling of labels in step 2 of the SW algorithm. Here again we need integrate posterior probability given the data. From theorem 2, it is possible to accept or reject a candidate label for a cluster based on the posterior probability $p(\pi'|data)$, where π' is a new partition of the graph after replacing a cluster's current label with new candidate label, so it is a *candidate partition*. Put it in another way, the acceptance probability is a function of $p(\pi'|data)$.

A comment on SW sampling: I believe the key difference is to replace prior with posterior term in step 1. That is, instead of using q_0 , we now use q_e which is a function of observed data. From the derivation in [Rubinstein and Kroese, 2008, chap 6.5], to use SW for sampling from posterior, it suffice to replace prior energy with posterior energy when sampling u , and the cluster label of x can just be sampled from the uniform distribution, just like the original SW algorithm. However, Barbu and Zhu uses a different sampling method for x in step 2, because that will further increase the convergence speed. Now the sampler in step 2 will choose those *good* candidate label in the posterior sense. The SWC-3 method also serves this purpose.

Question 2 (BOLD signal):

Information is transferred in axons between neurons by the release of neurotransmitter molecules at synapses. The interactions between neurotransmitter and receptors consume energy. Because energy is produced by oxidative metabolism, increased synaptic activity will also increase local demand for delivery of oxygen. This, counter-intuitively, increases the local blood flow, and increases the T_2^* image intensity [Lewin, 2003, Stroman, 2011]. Even deoxygenated hemoglobin decreases during neural activity, MR signal increases. This is because more oxygen is supplied to the brain region than is consumed.

Haemoglobin has different magnetic properties when it is bound to oxygen. This changes the local distortions of a magnetic field, which can be detected by MRI scanner. The relaxation times depend on the level of blood oxygenation, and the MRI signal depends on the relaxation time.

To understand the response of BOLD signal to a general stimulus signal, we need to know its response to a very short impulse signal, namely, a δ function. It is noted that the BOLD signal has about 2 seconds lag after the onset of stimulus. This initial dip is attributed to the transient increase of deoxygenated hemoglobin. After 4 to 6 seconds, the demand due to increased neuro activity results in an increased inflow of oxygenated blood, and it reaches the highest point. After the neuro activity stops, the BOLD goes below the baseline level. It takes about 20 seconds for the BOLD to go back to the baseline in this *poststimulus undershoot*. The balloon model [Buxton et al., 1998] is used to explain this extended period. The response function to a δ function is called the *Hemodynamic response function* (HRF). The response of a general boxcar function, would be the convolution of the boxcar function and the HRF.

Artificial signal in BOLD: Since BOLD signal does not measure neural activity directly, there might be some unwanted signal source in BOLD. First, changes in vascular system in response to the neural activity can occur in brain regions far from the actual activity location, initiated by flow-substances by neurons into the extracellular space, or initiated by glial cells (which do not directly participate in neural activity) [Huettel et al., Noll et al., 2001].

Another artificial signal source is the blood vessel. Because oxygen-rich blood flow increases dramatically to compensate the consumption of oxygen in neuronal activity, only a small proportion of oxygen will be actually used. The remaining hemoglobin-rich blood passes along the venous system, will increase the BOLD signal along the downstream of the activated neurons. Such effect leads to the question *Brain or Vein*, since may reported activated regions might be a sequence of venous drainage instead of a local neuronal activation.

A simple method to detect such signal confounds is the large magnitude change of the BOLD signal. Because of the large volume of blood vessel, BOLD signal will change in a larger magnitude due to the blood drainage. Another indicator is that the initial dip of the BOLD signal will not be observed in a blood vessel driven signal change. This is because the initial dip only happens when there is actual oxygen consumption happening in the capillaries, and will not happen at voxels containing large vessels that are distant from the active neurons.

Electric noise is a thermal driven Brownian movement, and can happen from the MRI coil, the

electronic hardware of the scanner. Thermal noise will be increased in a higher temperature environment. Ideally, the magnitude of thermal noise does not depend on the voxel's spatial location. One way to increase signal-to-noise (SNR) is to increase the B_0 . This is because the MR signal increase as a function of B_0^2 , while the noise increase as a function of B_0 . So th SNR is roughly a linear function of B_0 . Another method is to take the measurements multiple times so the noise cancels, but this does not apply to fMRI.

System noise comes from the variations of the scanner hardware. One example is the bias filed in MRI image. Another example is the scanner drift of fMRI. There is a linear trend on time courses. The linear trend can be subtracted by a detrending procedure by estimating the trend parameter.

The spatial resolution of fMRI is usually a few millimeters. If the spatial resolution is too low, each voxel will include different tissues. This is the *partial volume effect*. If a voxel includes 50% gray matter, 25% white matter, and 25% cerebrospinal fluid, the joint effect of these tissues would result in a intermediate value on the final voxel intensity.

Subject head motion is another source of unrelated signals in BOLD time course. With head motion, voxels at different time points may not correspond to the same spatial coordinates. Also it may introduce spatial correlations since all voxels will change at same time together. It can usually be estimated and removed by a motion correction method, i.e. register volumes at all time points to the first (or the middle) volume by a rigid-body transformation. Sometimes the motion is correlated with paradigm task signal and this makes motion correction potential dangerous since it may remove task-related signal also. Another method to mitigate motion is to do a linear regression with head motion parameters as explanatory variables and keep the residual signals for the following functional connectivity analysis.

Respiration is also one of the physiological noise in the BOLD signal. One solution is the band-pass filter that removes the signal below 0.01 Hz and above 0.1 Hz. However, if the TR is long, the breathing signal may be aliased and difficult to remove. Cardiac signal due to the hear beat may happen throughout the brain, but without a heart beat signal recorded together with fMRI scan, it is hard to remove such confounds.

There are other patterns of signals in BOLD, such as intrinsic variation in the brain due to practice, or due to fatigue, unrelated neural activity, cognitive or behavioral variability of how the subjects performs an task, etc. The variation between subjects can be alleviated by a group analysis with enough number of subjects.

Question 3 (MMSE estimator):

This is a Bayesian linear model, which is a Bayesian equivalent of the general linear model. The model can be written in a matrix notation as

$$\mathbf{x} = \mathbf{H}\boldsymbol{\theta} + \mathbf{w},$$

where \mathbf{x} is a $N \times 1$ data vector, $\mathbf{H} = [\mathbf{h}, \mathbf{1}]$ is $N \times 2$ matrix, and $\boldsymbol{\theta} = [\alpha, \beta]^\top$ is the parameter vector to be estimated.

(a) The MMSE estimate of $\boldsymbol{\theta}$ is the posterior mean $p(\boldsymbol{\theta}|\mathbf{x})$. This can be shown by compute the derivative of the Bayesian MSE

$$Bmse(\hat{\boldsymbol{\theta}}) = \mathbb{E}[(\boldsymbol{\theta} - \hat{\boldsymbol{\theta}})^2]$$

with respect to $\hat{\boldsymbol{\theta}}$. And the expectation is for both variable \mathbf{x} and $\boldsymbol{\theta}$ since both are assumed to be random in Bayesian framework. Therefore the MMSE estimator will be

$$\hat{\boldsymbol{\theta}} = \mathbb{E}[\boldsymbol{\theta}|\mathbf{x}]$$

Because both $p(\boldsymbol{\theta})$ and $p(\mathbf{x}|\boldsymbol{\theta})$ is Gaussian distribution, the joint probability $p(\mathbf{x}, \boldsymbol{\theta})$ is also Gaussian, and also is the posterior $p(\boldsymbol{\theta}|\mathbf{x})$. Here we directly use the result [Kay, ch. 10.6]

$$\mathbb{E}(\boldsymbol{\theta}|\mathbf{x}) = \boldsymbol{\mu}_\theta + \mathbf{C}_\theta \mathbf{H}^\top (\mathbf{H} \mathbf{C}_\theta \mathbf{H}^\top + \mathbf{C}_w)^{-1} (\mathbf{x} - \mathbf{H} \boldsymbol{\mu}_\theta),$$

where \mathbf{C}_θ is the covariance matrix of parameter vector $\boldsymbol{\theta}$. In this model, $\mathbf{C}_\theta = \sigma_\theta^2 \mathbf{I}$. \mathbf{C}_w is the covariance matrix of noise term \mathbf{w} . In our model, $\mathbf{C}_w = \mathbf{I}$, the $N \times N$ identity matrix. Plugging into the above equation, we get

$$\mathbb{E}(\boldsymbol{\theta}|\mathbf{x}) = \boldsymbol{\mu}_\theta + \sigma_\theta^2 \mathbf{I} \cdot \mathbf{H}^\top (\mathbf{H} \cdot \sigma_\theta^2 \mathbf{I}_{22} \cdot \mathbf{H}^\top + \mathbf{I})^{-1} (\mathbf{x} - \mathbf{H} \boldsymbol{\mu}_\theta),$$

where \mathbf{I}_{22} is 2×2 identity matrix. Using the identity $(\mathbf{A} + \mathbf{BCD})^{-1} = \mathbf{A}^{-1} - \mathbf{A}^{-1} \mathbf{B} (\mathbf{D} \mathbf{A}^{-1} \mathbf{B} + \mathbf{C}^{-1})^{-1} \mathbf{D} \mathbf{A}^{-1}$, we have

$$\begin{aligned} (\mathbf{H} \cdot \sigma_\theta^2 \mathbf{I}_{22} \cdot \mathbf{H}^\top + \mathbf{I})^{-1} &= \mathbf{I} - \mathbf{I} \mathbf{H} (\mathbf{H}^\top \mathbf{I} \mathbf{H} + \frac{1}{\sigma_\theta^2} \mathbf{I}_{22})^{-1} \mathbf{H}^\top \mathbf{I} \\ &= \mathbf{I} - \mathbf{H} (\mathbf{H}^\top \mathbf{H} + \frac{1}{\sigma_\theta^2} \mathbf{I}_{22})^{-1} \mathbf{H}^\top \end{aligned}$$

Because $\mathbf{H}^\top \mathbf{H} = \begin{pmatrix} N & 0 \\ 0 & N \end{pmatrix}$, we get

$$\mathbb{E}(\boldsymbol{\theta}|\mathbf{x}) = \boldsymbol{\mu}_\theta + \sigma_\theta^2 \mathbf{H}^\top (\mathbf{I} - \frac{\sigma_\theta^2}{1 + N\sigma_\theta^2} \cdot \mathbf{H} \mathbf{H}^\top) (\mathbf{x} - \mathbf{H} \boldsymbol{\mu}_\theta).$$

With some simplification, we have

$$\mathbb{E}(\boldsymbol{\theta}|\mathbf{x}) = \boldsymbol{\mu}_\theta + \frac{N}{N + \frac{1}{\sigma_\theta^2}} \left(\frac{1}{N} \mathbf{H}^\top \mathbf{x} - \boldsymbol{\mu}_\theta \right), \quad (3)$$

which is the MMSE estimator of the parameter vector $\boldsymbol{\theta}$.

(b). If we look at (3), we see it is a trade-off between the prior mean μ_{θ} , and the maximum likelihood (ML) estimation $\frac{1}{N}\mathbf{H}^{\top}\mathbf{x}$. If the prior variance σ_{θ}^2 is large (flat prior), the Bayesian MMSE estimator will converge to the ML estimation. It is called the likelihood *swamps out* the prior. Another condition to make this happen is when N is large. Again we trust data more than the prior.

(c). The MMSE minimize the quadratic cost function, and there are other definitions of cost function. If we define a Hit-or-miss cost function, the posterior mode will minimize such a function. Because in our model the posterior probability is Gaussian distribution, the posterior mean and posterior mode are identical. Therefore, the MAP estimator is also the MMSE estimator.

Question 4 (Fiber tracking):

(a). Diffusion Tensor Imaging (DTI): Water diffusion is anisotropic in white matter. This anisotropy can be measured by Diffusion Tensor Imaging and used to find the path of fiber tracts, and the connectivity between two region of interests (ROIs). The anisotropy can be represented by a diffusion ellipsoid. In DTI, The eigenvector of the diffusion tensor with the largest eigenvalue is assumed to align with the fiber orientation. For 3 eigenvalues $\lambda_1, \lambda_2, \lambda_3$, if at a specific voxel $\lambda_1 \gg \lambda_2 = \lambda_3$, the ellipsoid would be cylindrical structure along the direction of λ_1 . A straightforward approach for fiber tracking is connecting each voxel to the adjacent voxel along the largest eigenvalue's direction. [Mori et al. \[1999\]](#) propose to track a continuous rather than a discrete vector field, and is the first DTI fiber tracking method. The termination of fiber is judged by thresholding the fractional anisotropy value.

Flow based fiber tracking approach [[Campbell et al., 2005](#)] uses the full diffusion tensor as opposed to the principal diffusion directions, works better especially at in regions of partial volume. Still, the small noise of the diffusion tensors at each voxel can accumulate after hundreds or thousands of steps, resulting the estimated fiber deviating from the truth. Probabilistic tracking methods [[Basser et al., 2000](#)] are proposed to avoid this issue. At each step, the direction is sampled from a posterior distribution, the tracking is repeated to get multiple trajectories. A normalization step is followed to get the spatial distribution of the fiber bundle between two ROIs.

(b). The above methods fall into the category of *greedy* algorithms in that the tracking algorithm looks for a local optimal direction at each single step. [Zalesky \[2008\]](#) proposes a global optimization method that use dynamic programming technique to find the global optimal fiber path between two ROIs or seeds. Similar dynamic programming method are also used in [Poynton et al. \[2005\]](#), but no detailed information is available. Here I will take [Zalesky](#) as an example for the application of dynamic programming in fiber tracking.

[Zalesky](#) project the original fiber tracking problem to a new problem of finding shortest (or minimal-cost) path on a undirected weighted graph. A graph is defined such that each voxel is a node, and an edge is added whenever two voxels are spatially adjacent. The weights of the edges are defined by the diffusion tensor. That is, each edge is assigned a Bayesian posterior probability to indicate the likelihood that the edge is aligned to the true fiber. Between two seeds (ROIs), each possible path is assigned a probability of the confidence that this path is the true fiber. This probability is defined as the product of the posterior probability at each edges that it travels through. The goal is to find a path with largest probability. This is a global objective function and can be solved by dynamic programming method.

(c). Dynamic Programming (DP) [[Kleinberg, 2006](#), [Cormen, 2001](#)] is an optimization method that combine the optimal solution of subproblem to obtain the optimal solution of the original problem. Whenever there is a subproblem and a optimal solution to it, we see the possible application of DP. According to [[Kleinberg, 2006](#)], there are a few elements in the problem that we can apply DP:

- The number of subproblems should be polynomial.

- The solution of the original problem should be easily compute from the solutions of the subproblem. The cost of the solution to the original problem is the sum of the cost of subproblems and the cost of making the choice of which subproblem to use for solving the original problem [Cormen, 2001].
- There is natural ordering of the subproblems and recurrence that allows one to get the solution of a subproblem from that of a smaller subproblem.

DP is similar to divide-and-conquer method in that both split the original problem into subproblems. However, in divide-and-conquer, the subproblems are independent, while in DP the subproblems have overlaps, i.e., they share smaller subproblems. In such case, divide-and-conquer will do redundant work on these overlapping smaller subproblems. DP instead solve the common subproblems only once and have a data structure to save these solutions. This is called *memorization*, which is a crucial component of DP. By memorization we may achieve a polynomial-time algorithm when divide-and-conquer could not.

(d). Here we take Zalesky [2008] as example for searching the path searching between two points u and v in a graph. Define the graph $\mathcal{G} = (\mathcal{I}, \mathcal{E})$, where vertex set \mathcal{I} includes all voxels in a DTI image, and edge set \mathcal{E} includes any edge e if it is connected by two adjacent voxels. Assuming we already compute the edge weights $p(e)$ for all $e \in \{\mathcal{E}_x\}_{x \in \mathcal{I}}$ from the DTI images, the goal is to compute a discrete approximation (x_0, x_1, \dots, x_N) of a path with maximum probability $s_{u,v}^*$.

Data: Undirected \mathcal{G} , edge probabilities $p(e)$. Seed point u and target point v .

Result: $s_{u,v}^*$

```

1  $d(x) := -\infty, \forall x \in \mathcal{I};$ 
2  $prev(x) := empty, \forall x \in \mathcal{I};$ 
3  $d(u) := 0;$ 
4  $\mathbf{Q} := \mathcal{I};$ 
5 while  $\mathbf{Q}$  is not empty do
6    $\mathbf{Q} := \mathbf{Q} \setminus x^*,$  where  $d(x^*) = \max_{x \in \mathbf{Q}} d(x);$ 
7   foreach  $e \in \mathcal{E}_{x^*}$  do
8     if  $d(x^*) + \log p(e) > d(x_e)$  then
9        $d(x_e) := d(x^*) + \log p(e);$ 
10       $prev(x_e) := x^*;$ 
11    end
12  end
13 end
14  $x := v, \quad s_{u,v}^* := (v);$ 
15 while  $x \neq u$  do
16    $x := prev(x)$  Add  $x$  to the head of  $s_{u,v}^*;$ 
17 end
```

Algorithm 1: Dijkstra's algorithm for shortest path (maximum probability) search.

Algorithm 1 gives the Dijkstra's algorithm for shortest path search. Here we have minor

modification to apply it for maximum probability search for path between source u and target v . In this algorithm $d(x)$ is the maximum probability between u and x . $d(x)$ is initialized to $-\infty$. $prev(x)$ is the previous node of x along the path from u to v , and is initialized to empty. The algorithm also maintain a subset of vertex \mathbf{Q} for those that have not been explored (i.e. have not been computed for $d(x)$). At the beginning, \mathbf{Q} includes all nodes in the graph. After the node has been explored, it is removed from \mathbf{Q} (line 6). \mathbf{Q} uses a special data structure (priority queue) that can extract minimal value of $d(x)$ fast. When a x with minimal value of $d(x)$ is extracted, its adjacent neighbors x_e are explored by updating $d(x_e)$ (line 8). This is done until \mathbf{Q} is empty, that is, all nodes are explored.

A shortest path from u to v can be easily found by back searching the previous node of x at line 16 until the head node reaches source node u .

There are few things that is worth noting. First, standard Dijkstra’s algorithm only applies to graph with positive weight on all the edges, while the Bellman-Ford algorithm can be used to general with negative edge weights, but with a slower running time [Cormen, 2001, ch 24]. In our graph of DTI image, the edge weights $\log p(e)$ is negative. However, this is not a problem for Dijkstra’s algorithm since here we look for maximum probability of a path with all negative edge weights, which is equivalent to finding the minimal cost or shortest path with all positive edge weights.

Another thing is the implementation in 1 is not yet a dynamic programming approach (algorithm the author claims so), since it does not use DP concept, i.e. the memorization of the overlapping part of subproblems. This is not a problem as DP does not necessarily outperform greedy algorithm for all problems. Algorithm 1 is similar to the Breath-First-Search which is a greed algorithm. Not all greedy algorithm can achieve the optimal solution but Dijkstra’s algorithm can achieve it with a polynomial time, and is even faster than the dynamic programming version of the path-searching algorithm below.

The dynamic variants of shortest path search is due to Bellman-Ford algorithm. To give a sketch of the algorithm, think again splitting the optimal solution into subproblems. Assume $OPT(i,v)$ si the minimum cost of a path from i to v , the recursive formula used for DP can be written as [Kleinberg, 2006, chap 6.8]

$$\cdot \text{ if } i > 0 \text{ then } OPT(i, v) = \min(OPT(i - 1, v), \min_{w \in V}(OPT(i - 1, w) + c_{vw})),$$

where V is vertex set, and c_{vw} is the cost of edge (v, w) . If we assume the number of vertex (voxels) in the graph is n , and the number of edges is m , constructing a dynamic programming table in $\mathcal{O}(n^3)$ time, so the DP shortest path finding algorithm takes $\mathcal{O}(n^3)$ time to finish. A basic improvement of the algorithm is possible when the edges are close to n^2 (our DTI image even has much less edges), and have a running time $\mathcal{O}(nm)$. This seems still worse than the best Dijkstra algorithm with careful implementation of the priority queue on the \mathbf{Q} .

To see the difference of the path search between a 2D image and 3D volume, we see there is really no difference except of the increasing number of data points. To choose from different variants of the Dijkstra’s algorithm, one may need to know the number of edges compared to the number of vertex. For 2D regular lattice image, each pixel can be seen to have 4




Figure 1: Compare the speedup of path search on GPU (SSSP) with CPU. BFS is the standard breath-first-search algorithm. Figure is from [Harish and Narayanan \[2007\]](#).

neighbors. We will have $m = 2n$ in such case. For 3D volume with standard 6-neighborship, $m = 3n$. Either way m is the linear function of n , so the graph can be seen as sparse.

(e). How fast is the Dijkstra’s algorithm? The algorithm loops all n nodes. For each notes, it will reach all of its adjacent nodes to compute the minimal(maximal) value at line 8 of algorithm 1. For a graph with m edges, computing these minima will take $\mathcal{O}(m)$ time, so it will take $\mathcal{O}(nm)$ time at most to finish the Dijkstra’s algorithm.

For regular lattice in DTI volume, $m \approx 3n$, so it is a sufficiently sparse graph, and the upper bound of the running time is tighter than $\mathcal{O}(nm)$. The running time of Dijkstra’s algorithm depends on the implementation of the min-priority queue of \mathbf{Q} , and can be as low as $\mathcal{O}(m)$ time plus the time for n `extractMin` and m `ChangeKey` operations [[Kleinberg, 2006](#), chap 4.4]. If we implement the queue as a binary min-heap, the running time will be $\mathcal{O}(m \log n)$ [[Cormen, 2001](#), chap 24.3].

(f). Parallel implementation of Dijkstra’s algorithm has been available since [Nepomniaschaya and Dvoskina \[2000\]](#), but it is [Harish and Narayanan \[2007\]](#) that gives a implementation of the algorithm on Nvidia GPU platform. Two GPU kernel function is implemented to update the current short path data structure (i.e. the $d(x)$ in algorithm 1) to avoid the read after write inconsistencies, the speedup is shown in figure 1. It can be seen that the GPU version of the path-searching implementation has significant speedup compared to the CPU version, which means the fiber tracking on GPU is a promising research direction. Even more speed up can be shown in a newer implementation of [Buluē et al. \[2010\]](#), who has extensively use of matrix multiplication and thus has a high ratio of computation to communication.

(g). Before answer the question, we need a classification of fiber tracking algorithms. These algorithms are usually *single-path fiber tracking algorithms* and *all-paths fiber tracking problems*. The former find shortest (or most probable) path between know seed point u and end point v . The latter finds shortest paths starting from seed point u and ending to unknown target point. All-path approaches are often used to find a *probabilistic* path beginning from

a seed point.

The Dijkstra's algorithm can be used to find the all-paths as well as single path by same routine. This is because as a breath-first-search algorithm, finding the all-paths need the same amount of computation with finding the single-path.

Now for all-path searching, the seed and target point are both on surface of the 3D volume. I believe this does not have effect on the computation time of Dijkstra's algorithm. this is because as a breath-first-search algorithm, Dijkstra's algorithm may still have to explore all data points in the volume, even at last we only choose those shortest paths with ending point on the surface. If we apply algorithm 1 to all-path searching, the constraints that the ending points is on the surface only have effect on the set of v from which we want to back reconstruct the path. Because the back-searching of line 14 is linear time, changing of this piece does not change the upper bound of the Dijkstra's algorithm.

If, on the other hand, we want to find the all-pair shortest path between each pair of voxels in the volume, the constraint of starting and ending points on surface have larger impact on the running time. Without the constraint, We can run Dijkstra's algorithm for each $u \in \mathcal{I}$. Because Dijkstra's algorithm with binary min-heap implementation is $\mathcal{O}(m \log n)$, the running time for all-path shortest path searching will be $\mathcal{O}(mn \log n)$. With the constraints that $u \in \mathcal{I}_s$, where \mathcal{I}_s is the set of the points on the surface, the running time of all-pair searching will be $\mathcal{O}(|\mathcal{I}_s|m \log n)$, where $|\mathcal{I}_s|$ is the number of points on the surface.

A side note: I am thinking whether the global energy function is really optimal in the application sense. Suppose that we are looking for the global-optimal path from u to v . If the fiber tract we find pass through point x and y , and the probability of the sub-path between x and y is very low. We may argue even this path (u, v) maximize the global probability, the *broken* segment between x and y make me think the path (u, v) is not the optimal path in the real application. So, we may need a constraint: we do not want a path such that any segment of which has probability below certain threshold. If that happens, the fiber tract is believed to be broken and will not be the optimal solution. Finding the shortest path with constraints has been discussed in the literatures, and can be used for such problems.

Question 5 (Gibbs Sampling Convergence):

(a) If the starting point of the Markov chain is poorly chosen, the burn-in period will increase dramatically. The rule of thumb is choosing the starting sample close to the center of the distribution, i.e., the mode of the pdf. The proposal distribution of the Metropolis-Hasting sampling also has big impact on the steps needed to reach stationary distribution. For example, the random walker, a special case of Metropolis-Hasting sampler, has a symmetric proposal distribution (either uniform, or Normal distribution) with a tunable variance parameter. Increasing this variance parameter will have larger movement which is good to explore the whole support space, but at the risk of low acceptance rate and high correlation between samples. If the variance is too small, there are higher probability of accepting the candidates, but less opportunity to explore all modes of the target distribution, and the samples are also highly correlated. In such case, the chain will converge too slowly.

(b,c). Perhaps the single most popular approach is due to [Gelman and Rubin \[1992\]](#). To use their method, we need to run multiple parallel MCMC chains with different starting points. These chains must be overdispersed initially with respect to the posterior. Each chain has length $2N$ and the first half of points are discarded. If we use ψ_{mn} to represent the statistics of chain m at time n . The Gelman-Rubin method compute the between and within-sequence variances of the statistics ψ

$$\begin{aligned}
 B &= \frac{N}{M-1} \sum_{m=1}^M (\bar{\psi}_m - \bar{\psi})^2 \\
 \bar{\psi}_m &= \frac{1}{N} \sum_{n=1}^N \psi_{mn}, \quad \bar{\psi} = \frac{1}{M} \sum_{m=1}^M \bar{\psi}_m \\
 s_m^2 &= \frac{1}{N-1} \sum_{n=1}^N (\psi_{mn} - \bar{\psi}_m)^2 \\
 W &= \frac{1}{M} \sum_{m=1}^M s_m^2
 \end{aligned}$$

We can estimate the marginal posterior variance of the statistics ψ by a weighted average of W and B

$$\widehat{Var}(\psi|data) = \frac{N-1}{N} W + \frac{1}{N} B \quad (4)$$

This estimator overestimates the marginal variance when the starting points of the chain are overdispersed, but is unbiased when $n \rightarrow \inf$. On the other hand, the within-variance W will under-estimate the variance of ψ and converge to $Var(\psi)$ when $n \rightarrow \inf$. So we can compare the value of (4) with W . If they are very different, that means the chain is not converged yet. The Gelman-Rubin uses the estimated potential scale reduction

$$\hat{R} = \sqrt{\frac{\widehat{Var}(\psi|data)}{W}},$$

which declines to 1 as $n \rightarrow \inf$.

One issue with this method is to get overdispersed starting points, one need have some knowledge of the pdf of interest, for example, the modes and shape of high density regions. If multiple chains all start from a single mode of the density function, they may take long steps (if ever possible) to explore other modes. In such case, multiple chains do no help much compared to a single long chain, and the Gelman-Rubin method can not verify the convergence to the stationary distribution.

The Gelman-Rubin method is difficult to apply to sampler high dimensional random vectors, because saving multiple independent chains will require large memory, and sampling these chains also have high computation cost.

The second test for convergence is a non-parametric test. It is applied to the single chain. Assume $\theta^{(t)}$ is the statistics derived from the chain of $1, \dots, t$. When the chain reaches stationary, $\theta^{(t_1)}$ and $\theta^{(t_2)}$ has same marginal distribution for any t_1 and t_2 . Given a MCMC chain $\theta^{(1)}, \dots, \theta^{(T)}$, we can compare the empirical distributions of two half chain $(\theta_1^{(1)}, \dots, \theta_1^{(T/2)})$ and $(\theta_2^{(T/2)}, \dots, \theta_2^{(T)})$. The Kolmogorov-Smirnov statistics is defined as the supremum of the absolute value on the difference of two empirical distribution functions

$$K = \sup_{\eta} |F_1(\eta) - F_2(\eta)| = \frac{1}{M} \sup_{\eta} \left| \sum_{m=1}^M \mathbf{1}_1(\eta) - \sum_{m=1}^M \mathbf{1}_2(\eta) \right|,$$

where F_1 and F_2 is θ 's empirical distribution for two half chains, and $\mathbf{1}$ is indicator function. It is noted that because of the correlation between adjacent samples in MCMC, the half chain θ is sampled in a batch mode, i.e. θ_1^m and θ_1^{m+1} are separated by a interval to make a quasi-independent chain.

Under the stationary assumption, the limiting distribution of $\sqrt{M}K$ has the cdf [Robert and Casella, 2004, chap 12.2]

$$R(x) = 1 - \sum_{k=1}^{\infty} (-1)^{k-1} \exp\{-2k^2 x^2\}.$$

Now we can construct a hypothesis, with the null hypothesis as the two chains are from same distribution (i.e., the MCMC chain reaches stationary). The hull hypothesis is rejected if $\sqrt{M}K > K_{\alpha}$, where K_{α} is computed from $\Pr(x < K_{\alpha}) = 1 - \alpha$.

It is not straightforward to generalize the Kolmogorov-Smirnov test into higher dimension, especially for Ising model with dimension as the number of image voxels. From Wikipedia, this is because the maximum difference between two joint cumulative distribution functions is not generally the same as the maximum difference of any of the complementary distribution functions. One solution is to compare the cdfs of the two samples with all possible orderings, and take the largest of the set of resulting K-S statistics. In d dimensions, there are 2^{d-1} such orderings, which is intractable for Ising model.

(d). To test the convergence of a Gibbs sampler on a simple MRF such as a Ising model, we can instead compute the upper bound of the number of samplings. One method is to use the coupled sampled paths to study the convergence property [Johnson, 1996]. Two coupled process have same transition kernel but different starting point. The coupled paths reach

the same state after certain number of iterations. The iteration is defined as the sweep of all the data points in the mode. By examining the distribution of the iterations need for coupling, convergence properties of the sampling can be established. [Johnson](#) use a 64×64 regular lattice and assume a Ising model on the binary variables on the lattice. He tried to look for the relationship between the number of required sampling iterations and the Ising parameter β . [Figure 2](#) shows when β is small the required number of iterations is also small. However, because the growth of the iteration number is super-exponential, large value of β will need much more iterations to converge. When $\beta = 0.9$, the 95th and 99th quantile of the iterations distribution reached 1 million.

[Gibbs \[2000\]](#) also gives a upper bound of the iterations of Gibbs sampling on a one dimension Ising model. His upper bounds, however, is also a function of the square of data points, and of a tolerance ε , on the variation distance. Similar to [Johnson](#), the upper bound increase fast with β . For $\varepsilon = 0.01$, $\beta = 0.5$ gives a upper bound $128N^2$, while $\beta = 1.5$ gives $6162N^2$, where N is number of points in the lattice. [Gibbs](#) noted that there is no phase transition in this 1-D model. For higher dimensional Ising model, his upper bounds only applies to small β , and Convergence is slow when $\beta > \beta_0$, where β_0 is the reciprocal of the critical temperature. For a higher dimension of Ising model such that the number of neighbors of each data points increases, the convergence upper bound even increase faster. Actually the upper bound is a function of n , the number of neighbors of each node, and will increase when n increases.

As a side note, [\[Barbu and Zhu, 2005\]](#) talked about the convergence rate of Swendsen-Wang (S-W) algorithm on a Potts model. This is related to question 1 in this Qualify exam. They cited [Huber \[2003\]](#) who developed a new bounding chain algorithm that can diagnose the convergence of Swendsen-Wang sampling. The number of steps to reach perfect sampling (which means convergence to stationary distribution) is in the order of $\mathcal{O}(\log |E|)$, where E is the set of edges. This running time applies when the temperature is far below or far above critical temperature. [Cooper and Frieze \[1999\]](#) shows the mixing time (or convergence time) is polynomial if the number of neighbors of each node does not increase with $|V|$, the size of the nodes. This is good for regular lattice where the number of adjacent nodes is fixed, regardless of image size. Compared with the super-exponential rate of increase for the iteration number in standard Gibbs sampling, S-W algorithm is a big improvement for convergence rate. One thing that need note is these theoretical analysis is for the cases without external fields (data likelihood term).

References

- A. Barbu and S.C. Zhu. Generalizing swendsen-wang to sampling arbitrary posterior probabilities. *Pattern Analysis and Machine Intelligence, IEEE Transactions on*, 27(8):1239–1253, 2005.
- P.J. Basser, S. Pajevic, C. Pierpaoli, J. Duda, and A. Aldroubi. In vivo fiber tractography using dt-mri data. *Magnetic resonance in medicine*, 44(4):625–632, 2000.




Figure 2: Percentile of coupling iterations for Ising model of size 64×64 . Top curve shows the 99% and bottom shows the 95% percentile from the distribution of the iterations needed for coupling, as a function of β parameter. The percentiles are estimated using 1000 repetitions of Gibbs sampling initialized with all white and all-black value. Figure from [Johnson \[1996\]](#).

A.I. Buluç, J.R. Gilbert, and C. Budak. Solving path problems on the gpu. *Parallel Computing*, 36(5-6):241–253, 2010.

R.B. Buxton, E.C. Wong, and L.R. Frank. Dynamics of blood flow and oxygenation changes during brain activation: the balloon model. *Magnetic Resonance in Medicine*, 39(6):855–864, 1998.

J.S.W. Campbell, K. Siddiqi, V.V. Rymar, A.F. Sadikot, and G.B. Pike. Flow-based fiber tracking with diffusion tensor and q-ball data: validation and comparison to principal diffusion direction techniques. *NeuroImage*, 27(4):725–736, 2005.

C. Cooper and A.M. Frieze. Mixing properties of the swendsen-wang process on classes of graphs. *Random Structures and Algorithms*, 15(3-4):242–261, 1999.

T.H. Cormen. *Introduction to algorithms*. The MIT press, 2001.

A. Gelman and D.B. Rubin. Inference from iterative simulation using multiple sequences. *Statistical science*, 7(4):457–472, 1992.

A.L. Gibbs. Bounding the convergence time of the gibbs sampler in bayesian image restoration. *Biometrika*, 87(4):749–766, 2000.

G. Grimmett. *The random-cluster model*, volume 333. Springer Verlag, 2006.

P. Harish and P. Narayanan. Accelerating large graph algorithms on the gpu using cuda. *High Performance Computing-HiPC 2007*, pages 197–208, 2007.

- M. Huber. A bounding chain for swendsen-wang. *Random Structures & Algorithms*, 22(1): 43–59, 2003.
- S.A. Huettel, A.W. Song, and G. McCarthy. *Functional magnetic resonance imaging*, volume 1.
- V.E. Johnson. Studying convergence of markov chain monte carlo algorithms using coupled sample paths. *Journal of the American Statistical Association*, pages 154–166, 1996.
- S.M. Kay. Fundamentals of statistical signal processing, volume i: Estimation theory.
- J. Kleinberg. *Algorithm design*. Pearson Education India, 2006.
- J.S. Lewin. Functional mri: An introduction to methods. *Journal of Magnetic Resonance Imaging*, 17(3):383–383, 2003.
- S. Mori, B.J. Crain, VP Chacko, and P. Van Zijl. Three-dimensional tracking of axonal projections in the brain by magnetic resonance imaging. *Annals of neurology*, 45(2):265–269, 1999.
- A.S. Nepomniaschaya and M.A. Dvoskina. A simple implementation of dijkstra’s shortest path algorithm on associative parallel processors. *Fundamenta Informaticae*, 43(1):227–243, 2000.
- D.C. Noll et al. A primer on mri and functional mri. *Technical Report, University of Michigan, Ann Arbor, USA*, 2001.
- C. Poynton, R. Lal, JT Ratnanather, S. Mori, D. Boatman, and MI Miller. Probabilistic tracking of fiber pathways using dynamic programming. In *11th Annu. Meeting Org. Human Brain Mapp., Toronto, ON, Canada*, 2005.
- C.P. Robert and G. Casella. *Monte Carlo statistical methods*. Springer Verlag, 2004.
- R.Y. Rubinstein and D.P. Kroese. *Simulation and the Monte Carlo method*, volume 707. Wiley-interscience, 2008.
- P.W. Stroman. *Essentials of Functional MRI*. CRC Press, 2011.
- G. Winkler. *Image analysis, random fields and Markov chain Monte Carlo methods: a mathematical introduction*, volume 27. Springer Verlag, 2003.
- A. Zalesky. Dt-mri fiber tracking: a shortest paths approach. *Medical Imaging, IEEE Transactions on*, 27(10):1458–1471, 2008.